

AIO Engineer Technical Assessment

Shift Scheduling System
Rust

Position	AIO Engineer
Level	Mid / Senior
Duration	7 days
Format	Independent, remote
Submission	Git repository

1. Overview

Design and implement a shift scheduling system composed of two microservices:

Service	Responsibility	Storage
Data Service	Staff and group management	PostgreSQL
Scheduling Service	Asynchronous shift schedule generation	PostgreSQL

Each service should be independently deployable and communicate over HTTP.

2. Technology Stack

Component	Requirement
Language	Rust (edition 2021+)
Web Framework	Axum
Async Runtime	Tokio
Database	PostgreSQL with sqlx
Serialization	serde / serde_json
Containerization	Docker, docker-compose
API Documentation	utoipa (OpenAPI / Swagger)
Caching	Redis

3. Data Service

3.1 Entities

- **Staff** - id, name, email (unique), position, status (ACTIVE / INACTIVE), timestamps
- **Staff Group** - id, name, parent group reference (supports nesting)
- **Group Membership** - many-to-many relationship between staff and groups

3.2 Database

Table	Purpose
staff	Staff records
staff_groups	Groups with self-referencing parent for hierarchy
group_memberships	Staff-to-group associations

Column types, constraints, and indexes are left to the candidate.

3.3 API

Candidates are expected to design a complete set of RESTful APIs, including:

- Full CRUD operations for staff and groups
- Adding and removing staff from groups
- Batch import from JSON files for all resources
- Hierarchical group resolution - retrieving all members under a group, including members of nested subgroups

Reference endpoint:

```
GET /api/v1/groups/{id}/resolved-members
```

3.4 Caching

Read-heavy endpoints should be cached using Redis.

4. Scheduling Service

4.1 Responsibility

- Accept schedule generation requests and process them asynchronously
- Retrieve staff data from the Data Service over HTTP
- Persist job state and generated schedules in PostgreSQL

4.2 Shift Types

Type	Description
MORNING	Morning shift
EVENING	Evening shift
DAY_OFF	Day off

4.3 API

Submit a schedule generation request

```
POST /api/v1/schedules
```

Request:

```
{
  "staff_group_id": "group-123",
  "period_begin_date": "2025-05-19"
}
```

Response (202 Accepted):

```
{
  "schedule_id": "job-abc123",
  "status": "PENDING"
}
```

Check job status

```
GET /api/v1/schedules/{schedule_id}/status
```

Status values: PENDING, PROCESSING, COMPLETED, FAILED

Retrieve generated schedule

```
GET /api/v1/schedules/{schedule_id}/result
```

Response:

```
{
  "schedule_id": "job-abc123",
  "period_begin_date": "2025-05-19",
  "staff_group_id": "group-123",
  "assignments": [
```

```
{
  { "staff_id": "staff-1", "date": "2025-05-19", "shift": "MORNING" },
  { "staff_id": "staff-1", "date": "2025-05-20", "shift": "EVENING" },
  { "staff_id": "staff-1", "date": "2025-05-21", "shift": "DAY_OFF" }
]
}
```

4.4 Scheduling Rules

Input

- A staff group, fetched from Data Service via the resolved-members endpoint
- A period start date (`period_begin_date`), which must fall on a Monday
- Period length: 28 days (4 weeks)

Constraints

1. Each staff member is assigned exactly one shift per day: MORNING, EVENING, or DAY_OFF.
2. Each staff member must have at least 1 and at most 2 DAY_OFF entries per week (Monday through Sunday).
3. A MORNING shift must not immediately follow an EVENING shift. If day N is assigned EVENING, then day N+1 must not be MORNING.
4. Only staff with status ACTIVE should be included in scheduling.
5. On any given day, the difference between the number of staff assigned to MORNING and the number assigned to EVENING must not exceed a configurable threshold. This ensures balanced daily coverage across shifts.

Example - valid assignment (1 staff, 1 week)

Mon	MORNING
Tue	EVENING
Wed	EVENING
Thu	DAY_OFF
Fri	MORNING
Sat	MORNING
Sun	EVENING

Total: 3 MORNING, 3 EVENING, 1 DAY_OFF

All constraints satisfied.

Example - valid daily coverage (4 staff, Monday)

Staff-1	MORNING
Staff-2	MORNING
Staff-3	EVENING
Staff-4	EVENING

MORNING: 2, EVENING: 2, difference: 0

Satisfies constraint 5.

Example - invalid daily coverage (4 staff, Monday, `max_daily_shift_diff = 1`)

Staff-1	MORNING
Staff-2	MORNING

```

Staff-3 MORNING
Staff-4 DAY_OFF

MORNING: 3, EVENING: 0, difference: 3
INVALID - violates constraint 5.

```

Example - invalid assignment

```

Mon EVENING
Tue MORNING      INVALID - violates constraint 3

```

Configuration

Each rule should be configurable (enable/disable and parameter adjustment) via a configuration file:

Rule	Key	Default
Minimum days off per week	min_day_off_per_week	1
Maximum days off per week	max_day_off_per_week	2
Disallow MORNING after EVENING	no_morning_after_evening	true
Maximum daily shift difference	max_daily_shift_diff	1

4.5 Database

Table	Purpose
schedule_jobs	Job metadata - group, period, status, timestamps
shift_assignments	Generated assignments - staff, date, shift type

Column types, constraints, and indexes are left to the candidate.

4.6 Error Handling

- Failed jobs must be logged and marked with status FAILED.
- Errors from the Data Service should be handled gracefully without crashing the process.

5. Architecture

- Apply **Clean Architecture** principles with clear separation between API, domain, and infrastructure layers.
- Use **traits** for dependency inversion between layers.
- Scheduling rules must be implemented as **configurable components** that can be enabled or disabled via configuration.
- Organize the project as a **Cargo workspace**:

```
shift-scheduler/
├── Cargo.toml          # workspace root
├── docker-compose.yml
├── data-service/
│   ├── Cargo.toml
│   └── src/
├── scheduling-service/
│   ├── Cargo.toml
│   └── src/
└── shared/
    ├── Cargo.toml
    └── src/
└── sample-data/
    ├── staff.json
    └── groups.json
```

6. Testing

- **Unit tests** covering scheduling rules and domain logic.
- **Integration tests** covering API endpoints.
- Use **mockall** to mock Data Service dependencies within the Scheduling Service.
- All code must pass `cargo clippy` with no warnings and be formatted with `cargo fmt`.

7. Deliverables

7.1 Source Code

Two services organized within a single Cargo workspace, including a shared crate for common types.

7.2 Docker Compose

A `docker-compose.yml` that includes the following containers:

Container	Description
data-service	Data Service instance
scheduling-service	Scheduling Service instance

postgresql	PostgreSQL database
redis	Redis cache

The entire system must start with a single command: `docker-compose up`.

7.3 Documentation

A `README.md` covering:

- Build and run instructions
- API documentation (Swagger UI link or Postman collection)
- Database schema or ERD

7.4 Sample Data

JSON files suitable for batch import (staff, groups).

8. Evaluation Criteria

Criteria	Weight
Architecture and design - Clean Architecture, trait usage, layer separation	25%
Rust code quality - idiomatic patterns, error handling, effective use of the type system	25%
Feature completeness - CRUD, batch import, scheduling, async processing	25%
Testing and DevOps - test coverage, Docker setup, documentation	25%

Bonus

The following are not required but will be considered favorably:

- Automatic cache invalidation on data mutations
- Type-state pattern for job lifecycle management
- Graceful shutdown handling
- Distributed tracing via OpenTelemetry
- Compile-time query verification with sqlx
- Circuit breaker pattern for Data Service calls
- Distributed locking to prevent duplicate job processing